


Analysis:


1. Word frequency analysis:

Word Count Histogram:




As can be seen, most words in the vocabulary have a count of less than 100. Let's look at the histogram for counts for various ranges

when counts < 100:




```
len(counts.loc[counts == 1])  
3266
```

when counts > 100 and counts <=10000



when counts > 10000



let's look at the words with very high counts

```
[('the', (0, 833516)), ('and', (1, 678137)), ('to', (2, 555648)), ('a', (3, 543555)), ('was', (4, 391267)), ('she', (5, 262987)), ('he', (6, 254659)), ('it', (7, 251852)), ('they', (8, 246987)), ('her', (9, 225202)), ('lily', (10, 193492)), ('day', (11, 157142)), ('with', (12, 149212)), ('said', (13, 148892)), ('his', (14, 147392)), ('in', (15, 135190)), ('that', (16, 133920)), ('but', (17, 116605)), ('s', (18, 111434)), ('you', (19, 110824))]
```

These are mostly articles, prepositions, verbs, adjectives and adverbs, used frequently.

Training Procedure:

Dataset:

Looking at the word count distribution and following the method described in paper, we have used the entire dataset for training. Few reasons are highlighted below:

1. Ideally the quality of embeddings is checked using downstream tasks. Google provided 20000 syntactic, semantic similarity tasks for evaluation.

2. We can create a held out set from the training corpus for evaluation using the same loss function however, since a large amount of words (3266) occur only once in corpus, hence random splitting of corpus leads to absence of those words from training, which will result in no training for these words.
3. We train till the loss function flatlines
4. Vocab size = 14040
5. Since the corpus size is less we choose an embedding dimension of 64, to reduce overfitting(but has to be checked with downstream tasks). T value is also increased so that relatively more frequent words are selected for training. Negative sample size is 10, as directed in paper (they instruct to use between 5-20, for small datasets)

Important Parameters:

SKIPGRAM_N_WORDS = 7

T = 1e-3

NEG_SAMPLES = 10

EMBED_DIM = 64

EPOCHS = 10

LOSS = Binary Cross Entropy(as described in paper)

Training loss variation :

1028.1s	5	EPOCH: 1, Train Loss: 0.9059986357441109
1959.7s	6	EPOCH: 2, Train Loss: 0.3159135907266534
2889.5s	7	EPOCH: 3, Train Loss: 0.2891594208477971
3822.9s	8	EPOCH: 4, Train Loss: 0.28171491454300107
4756.6s	9	EPOCH: 5, Train Loss: 0.2783028142673246
5698.6s	10	EPOCH: 6, Train Loss: 0.2763740879335526
6628.4s	11	EPOCH: 7, Train Loss: 0.2751621067954487
7556.8s	12	EPOCH: 8, Train Loss: 0.27431692594414153
8495.7s	13	EPOCH: 9, Train Loss: 0.27369742488928056
9421.1s	14	EPOCH: 10, Train Loss: 0.27321662652167233

Results:

Since the dataset is extremely small, the embeddings are not able to extract all the semantic and syntactic meanings. But we have evaluated the model on some similarity and analogical tasks("Germany" : "Berlin" :: "France" : ? ans:"Paris"). Results are shown below:

```
model.get_similar_words('food', 10)
```

```
['quiz',  
 'cheese',  
 'animal',  
 'disappoint',  
 'snack',  
 'dumped',  
 'apples',  
 'habits',  
 'neatly',  
 'cook']
```

```
model.get_similar_words('chocolate', 10)
```

```
['cocoa',  
 'drink',  
 'ingredients',  
 'butter',  
 'cookies',  
 'soup',  
 'treat',  
 'bowl',  
 'wound',  
 'sugar']
```

```
model.get_similar_words('love', 10)
```

```
['welcome',  
'thank',  
'm',  
'are',  
'have',  
'will',  
'you',  
'sorry',  
'smiles',  
'thanks']
```

```
model.get_similar_words('animal', 10)
```

```
['leopard',  
'cheerful',  
'zoo',  
'food',  
'mouse',  
'the',  
'min',  
'ollie',  
'lucy',  
'mimi']
```

```
model.get_similar_words('dolls', 10)
```

```
['dress-up',  
'bears',  
'cars',  
'with',  
'pictures',
```

```
model.get_similar_words('king', 10)
```

```
['castle',  
'near',  
'cage',  
'ollie',  
'cow',  
'poor',  
'protect',  
'jealous',  
'elephant',  
'value']
```

```
model.get_similar_words('lion', 10)
```

```
['brave',  
 'fox',  
 'tim',  
 'leo',  
 'rabbit',  
 'monkey',  
 'bobo',  
 'freddy',  
 'monster',  
 'frog']
```

```
model.get_analogy('prince', 'princess', 'boy', 10)
```

```
['friends',  
 'ducky',  
 'above',  
 'brave',  
 'heel',  
 'pigeon',  
 'task',  
 'problem',  
 'ending',  
 'little']
```

As can be seen some of the similar words do make sense, like a king will live in a castle, chocolate and cocoa are similar, children do make their dolls play dress-up etc. Results on analogy tasks are not that great in our observations hence only one example is shown.

We can do more experiments by changing the size of embeddings, and the size of negative sampling for better contrastive estimation.

Here is a 2d plot of some sampled word embeddings using T-SNE dimensionality reduction.

